



Gorgo Go for Java Programmers

May 31, 2026

Contents

14 Essential Standard Library	1
fmt — Revisited	1
io — The Glue of Go I/O	2
bufio — Buffered I/O	3
os — Files and the Process Environment	4
os/exec — Running Subprocesses	5
flag — CLI Flag Parsing	6
time — Clocks, Durations, and Timers	7
path/filepath — Cross-Platform Path Manipulation	8
log/slog — Structured Logging	9
regexp — Regular Expressions	11
cmp and maps — Comparison and Map Utilities	12
encoding/base64 — Base64 Encoding	14
crypto/rand and math/rand/v2 — Random Numbers	15
crypto/sha256, crypto/aes, crypto/cipher — Hashing and Encryption	15
Key Points	16
Exercises	17

Chapter 14

Essential Standard Library

Go's standard library covers most of what you reach for in daily backend work — I/O, file access, time, logging, CLI flags, and pattern matching — without pulling in external dependencies. This chapter walks through the packages every Go programmer uses constantly, with Java comparisons where the mental model transfer is non-obvious.

fmt — Revisited

Chapter 1 introduced `fmt.Println`, `fmt.Printf`, and `fmt.Sprintf`. This section adds the verbs and writer-directed functions you will use daily.

Diagnostic Verbs

Three verbs are especially useful for debugging:

Verb	Output
<code>%v</code>	Default format — values only
<code>%+v</code>	Struct with field names
<code> %#v</code>	Go-syntax representation; can paste back in code
<code>%T</code>	Go type name

```
type Track struct {
    Title string
    Artist string
    BPM   int
}

t := Track{Title: "positions", Artist: "Ariana Grande", BPM: 114}
fmt.Printf("%v\n", t)    // {positions Ariana Grande 114}
fmt.Printf("%+v\n", t)  // {Title:positions Artist:Ariana Grande BPM:114}
fmt.Printf("%#v\n", t)  // main.Track{Title:"positions", Artist:"Ariana Grande", BPM:114}
fmt.Printf("%T\n", t)   // main.Track
```

`%+v` is your first move when inspecting an unknown struct value. `%#v` gives you a snippet you can paste directly into a test or a var declaration.



Tip: `%#v` on a slice prints `[]int{1, 2, 3}` rather than `[1 2 3]`. It is more verbose but unambiguous — very useful in test failure messages.

fmt.Fprintf to Any io.Writer

```
func Fprintf(w io.Writer, format string, a ...any) (n int, err error)
```

`fmt.Printf(format, args...)` is simply `fmt.Fprintf(os.Stdout, format, args...)`. Passing any `io.Writer` redirects the output:

```
fmt.Fprintf(os.Stderr, "warn: retrying in %v\n", delay)
fmt.Fprintf(logFile, "[INFO] loaded %d tracks\n", n)
fmt.Fprintf(&buf, "SELECT * FROM tracks WHERE bpm > %d", minBPM)
```

Because `*os.File`, `*bytes.Buffer`, `*strings.Builder`, `net.Conn`, and `http.ResponseWriter` all satisfy `io.Writer`, a single `Fprintf` call works with any of them.



Tip: In Java you might have separate `PrintStream`, `PrintWriter`, and `StringBuilder.append` paths. In Go there is one path: accept an `io.Writer`, call `fmt.Fprintf`.

io — The Glue of Go I/O

The `io` package defines the core interfaces that tie all I/O in Go together. You saw `io.Reader` and `io.Writer` in Chapter 8; this section covers the functions that compose them.

io.ReadAll and io.Copy

```
func ReadAll(r Reader) ([]byte, error)           // read r until EOF; return all bytes
func Copy(dst Writer, src Reader) (int64, error) // copy src to dst until EOF or error
```

`io.ReadAll` is the Go equivalent of Java's `InputStream.readAllBytes()` (Java 9+). Use it when you need the entire contents in memory:

```
data, err := io.ReadAll(resp.Body)
```

`io.Copy` streams from a `Reader` to a `Writer` without loading everything into memory at once. It is Go's answer to Java's `InputStream.transferTo(OutputStream)`:

```
n, err := io.Copy(dst, src) // n is the number of bytes copied
```



Trap: `io.ReadAll` on a large HTTP response body allocates the entire response in memory. For large or unbounded responses, prefer `io.Copy` to stream directly to a file or another writer.

Composing Readers and Writers

The `io` package provides several functions that wrap and combine readers and writers without copying data:

```
func MultiReader(readers ...Reader) Reader           // reads from each reader in sequence
func MultiWriter(writers ...Writer) Writer           // writes to all writers simultaneously
func TeeReader(r Reader, w Writer) Reader            // returns reader that copies to w as it reads
func LimitReader(r Reader, n int64) Reader           // reads at most n bytes from r
func Pipe() (*PipeReader, *PipeWriter)              // creates synchronous in-memory pipe
```

MultiReader is like Java's `SequenceInputStream`. MultiWriter is like Java's `TeeOutputStream` but for any number of writers. TeeReader is useful for logging raw bytes while processing them:

```
// Read from r and simultaneously write everything read to log.
logged := io.TeeReader(r, log)
io.Copy(processor, logged)
```

LimitReader prevents reading past a byte budget:

```
limited := io.LimitReader(req.Body, 1<<20) // never read more than 1 MiB
data, err := io.ReadAll(limited)
```

io.Pipe creates a synchronous in-memory pipe: writes to the PipeWriter block until the PipeReader consumes them. It connects a producer goroutine to a consumer that expects an `io.Reader`, without a buffer:

```
pr, pw := io.Pipe()
go func() {
    fmt.Fprintln(pw, "thank u, next") // Ariana Grande
    pw.Close()
}()
data, _ := io.ReadAll(pr)
fmt.Println(string(data)) // thank u, next
```



Tip: The `io` composition functions never allocate a large intermediate buffer. They are building blocks for streaming pipelines: `LimitReader` guards against oversized input, `TeeReader` adds tap-style logging, and `MultiWriter` fans out to several sinks.

Other io Signatures

```
func ReadFull(r Reader, buf []byte) (n int, err error) // reads exactly len(buf) bytes
func WriteString(w Writer, s string) (n int, err error) // writes a string to w
var Discard io.Writer // discards all writes; useful in tests
```

`io.Discard` is a writer that throws data away — handy in tests when you want to drain a reader without storing the bytes.

bufio — Buffered I/O

Raw `io.Reader` and `io.Writer` operations may call the OS for every byte. `bufio` wraps any reader or writer in a userspace buffer, batching syscalls for performance. The Java equivalent is `BufferedReader` / `BufferedWriter`.

bufio.Scanner

Scanner is the idiomatic way to read text line by line (or word by word, or any custom token):

```
func NewScanner(r io.Reader) *Scanner

scanner := bufio.NewScanner(os.Stdin)
for scanner.Scan() { // Scan returns false at EOF or error
    line := scanner.Text() // current line, no trailing newline
    fmt.Println(line)
}
if err := scanner.Err(); err != nil {
    log.Fatal(err)
}
```

The loop pattern — for `scanner.Scan()` — is the Go equivalent of Java's `BufferedReader.readLine() != null`.



Trap: Always check `scanner.Err()` after the loop. `Scan` returns `false` both at EOF (no error) and on a read error. Skipping the check silently drops I/O errors. [*no-discard-error*]

By default `Scanner` splits on newlines. You can change the split function:

```
scanner.Split(bufio.ScanWords) // split on whitespace
scanner.Split(bufio.ScanBytes) // one byte at a time
scanner.Split(bufio.ScanRunes) // one UTF-8 rune at a time
```

`Scanner` has a default buffer of 64 KiB per token. For very long lines, call `scanner.Buffer(make([]byte, cap), cap)` before the first `Scan`.

bufio.NewReader and bufio.NewWriter

```
func NewReader(rd io.Reader) *Reader // wraps rd in a 4096-byte read buffer
func NewWriter(w io.Writer) *Writer  // wraps w in a 4096-byte write buffer
```

Use `bufio.NewReader` when you need `ReadString`, `ReadByte`, or `Peek` on an existing `io.Reader`:

```
br := bufio.NewReader(conn)
header, err := br.ReadString('\n') // read up to and including the newline
```

Use `bufio.NewWriter` to batch small writes to an expensive underlying writer (a file or network connection). You **must** call `Flush` when done or buffered data will be silently discarded:

```
bw := bufio.NewWriter(file)
fmt.Fprintln(bw, "Don't Start Now") // Dua Lipa --- written to buffer, not file yet
bw.Flush()                          // now the data reaches the file
```



Trap: Forgetting `bw.Flush()` is one of the most common Go I/O bugs. Use `defer bw.Flush()` immediately after creating the `bufio.Writer` so you cannot forget it.

os — Files and the Process Environment

Opening and Creating Files

```
func Open(name string) (*File, error) // open for reading only
func Create(name string) (*File, error) // create or truncate, read/write
func OpenFile(name string, flag int, perm FileMode) (*File, error) // full control over flags/permissions
func ReadFile(name string) ([]byte, error) // read entire file into memory
func WriteFile(name string, data []byte, perm FileMode) error // write entire file, creating if needed
```

For most file operations you need just two patterns:

Read the whole thing:

```
data, err := os.ReadFile("playlist.json")
if err != nil {
    return fmt.Errorf("reading playlist: %w", err)
}
```

Stream a large file:


```
f, err := os.Open("tracks.csv")
if err != nil {
    return err
}
defer f.Close()
scanner := bufio.NewScanner(f)
for scanner.Scan() {
    process(scanner.Text())
}
```

*os.File satisfies both io.Reader and io.Writer, so it plugs directly into any function that accepts those interfaces.



Trap: Always defer `f.Close()` immediately after a successful `os.Open` or `os.Create`. A file left open leaks an OS file descriptor. Go does not close files when they go out of scope the way Java's try-with-resources does.

`os.WriteFile` is the simplest way to atomically replace a small file:

```
err := os.WriteFile("config.json", data, 0o644)
```

Process Environment

```
var Args []string           // command-line arguments; Args[0] is the program name
var Stdin *File             // standard input
var Stdout *File            // standard output
var Stderr *File            // standard error
func Getenv(key string) string // returns the value of an environment variable
func LookupEnv(key string) (string, bool) // like Getenv but distinguishes missing from empty
```

`os.Stdin`, `os.Stdout`, and `os.Stderr` are ordinary `*os.File` values that satisfy `io.Reader` / `io.Writer`. That is why `fmt.Fprintln(os.Stderr, msg)` works without any special cast.

```
token := os.Getenv("API_TOKEN")
if token == "" {
    fmt.Fprintln(os.Stderr, "API_TOKEN not set")
    os.Exit(1)
}
```

`os.Args[0]` is the program name; `os.Args[1:]` are the user-supplied arguments — identical to Java's `String[] args` but global. For real CLI tools, use the `flag` package (see below) rather than parsing `os.Args` by hand.

os/exec — Running Subprocesses

The `os/exec` package runs external programs. The Java counterpart is `ProcessBuilder`.

```
func Command(name string, arg ...string) *Cmd // builds a Cmd ready to run

// Run a command and capture its combined output.
out, err := exec.Command("git", "rev-parse", "--short", "HEAD").Output()
if err != nil {
    log.Fatal(err)
}
fmt.Printf("commit: %s", out)
```

`Cmd.Output()` runs the command, waits for it to finish, and returns stdout as a []byte. Use `Cmd.Run()` when you do not need the output:

```
cmd := exec.Command("gofmt", "-w", "main.go")
cmd.Stdout = os.Stdout // forward subprocess stdout to our stdout
cmd.Stderr = os.Stderr // forward subprocess stderr to our stderr
if err := cmd.Run(); err != nil {
    log.Fatalf("gofmt failed: %v", err)
}
```

Piping stdio

`Cmd.Stdin`, `Cmd.Stdout`, and `Cmd.Stderr` are `io.Reader` / `io.Writer` fields. Assign any compatible value:

```
cmd := exec.Command("sort")
cmd.Stdin = strings.NewReader("banana\napple\ncherry\n")
cmd.Stdout = os.Stdout
cmd.Run()
// apple
// banana
// cherry
```

For a pipeline, connect one command's stdout to another's stdin using `io.Pipe` or `cmd.StdoutPipe()`:

```
c1 := exec.Command("cat", "tracks.txt")
c2 := exec.Command("grep", "pop")
c2.Stdin, _ = c1.StdoutPipe() // c2 reads from c1's stdout
c2.Stdout = os.Stdout
c2.Start()
c1.Run()
c2.Wait()
```



Trap: `exec.Command` does not invoke a shell. `exec.Command("ls -la")` does **not** work — it looks for a program literally named "ls -la". Pass each argument separately: `exec.Command("ls", "-la")`.

flag — CLI Flag Parsing

The `flag` package parses command-line flags in the style of Go tools themselves. The Java equivalent is Apache Commons CLI or picocli, but `flag` is built in and needs no dependency.

```
func Bool(name string, value bool, usage string) *bool // define a boolean flag
func Int(name string, value int, usage string) *int    // define an integer flag
func String(name string, value string, usage string) *string // define a string flag
func Parse()                                           // parse os.Args[1:]

var (
    port    = flag.Int("port", 8080, "HTTP port to listen on")
    verbose = flag.Bool("verbose", false, "enable verbose logging")
    output  = flag.String("output", "stdout", "output file path")
)

func main() {
    flag.Parse() // must call before reading flag values
    fmt.Printf("starting on :%d (verbose=%v)\n", *port, *verbose)
}
```

Running `./app -port 9090 -verbose` sets `*port` to 9090 and `*verbose` to true. Running `./app -help` prints the usage string automatically.

Non-flag arguments (positional arguments after all flags) are available via `flag.Args()`.



Tip: Call `flag.Parse()` in main, not in init. Calling it from init makes testing harder because init runs before your test code can set up the argument list.

For more complex sub-command CLIs, the third-party cobra library (used by kubectl, Docker, and the Go CLI itself) builds on top of flag idioms.

time — Clocks, Durations, and Timers

time.Duration and time.Time

`time.Duration` is an int64 counting nanoseconds. Named constants give you human-readable literals:

```
const (
    Nanosecond Duration = 1
    Microsecond      = 1000 * Nanosecond
    Millisecond       = 1000 * Microsecond
    Second            = 1000 * Millisecond
    Minute            = 60 * Second
    Hour              = 60 * Minute
)
```

You write durations as arithmetic:

```
timeout := 30 * time.Second // 30s
jitter  := 500 * time.Millisecond
```

In Java you use `Duration.ofSeconds(30)` and `TimeUnit.MILLISECONDS`. Go's arithmetic on named constants is more concise.

`time.Time` represents an instant in time with nanosecond precision. The zero value (`time.Time{}`) is January 1, year 1, UTC — not null.

```
now := time.Now() // current local time
fmt.Println(now.Format(time.RFC3339)) // 2024-11-22T15:04:05-08:00
fmt.Println(now.UTC().Format(time.RFC3339)) // same instant, UTC offset
```



Wut: Go's reference time for layout strings is Mon Jan 2 15:04:05 MST 2006 — a specific instant, not placeholder tokens like Java's yyyy-MM-dd. `time.RFC3339` is the constant "2006-01-02T15:04:05Z07:00". If your format string looks wrong, check that you used the reference digits (month=01, day=02, hour=15, minute=04, second=05, year=2006).

time.Now, time.Since, and time.After

```
func Now() Time // current time
func Since(t Time) Duration // elapsed time since t; equivalent to Now().Sub(t)
func After(d Duration) <-chan Time // returns a channel that receives after duration d
func Sleep(d Duration) // blocks the calling goroutine for d

start := time.Now()
doWork()
fmt.Printf("finished in %v\n", time.Since(start))
```

`time.After` is the idiomatic timeout in a `select`:

```
select {
case result := <-work:
    fmt.Println("got result:", result)
case <-time.After(5 * time.Second):
    fmt.Println("timed out")
}
```

`time.Ticker` and `time.Timer`

```
func NewTicker(d Duration) *Ticker // fires every d; Ticker.C is the receive channel
func NewTimer(d Duration) *Timer   // fires once after d; Timer.C is the receive channel
```

`time.Ticker` delivers a tick on its channel at every interval — the same concept as Java’s `ScheduledExecutorService.schedule`

```
ticker := time.NewTicker(1 * time.Second)
defer ticker.Stop() // stop the ticker when you are done or it leaks a goroutine
```

```
for range ticker.C {
    fmt.Println("tick:", time.Now().Format(time.TimeOnly))
}
```

`time.Timer` fires once after a delay. If you do not need the channel — you just want to pause — use `time.Sleep` instead.



Trap: Always call `ticker.Stop()` (or `timer.Stop()`) when you are done with a `Ticker` or `Timer`. Failing to stop them leaks the internal goroutine that drives the channel.

`path/filepath` — Cross-Platform Path Manipulation

`path/filepath` handles OS path separators correctly (`\` on Windows, `/` on Unix), unlike the `path` package which is hardcoded to forward slashes. The Java equivalent is `java.nio.file.Path`.

```
func Join(elem ...string) string // joins path elements with the OS separator
func Dir(path string) string      // returns all but the last element of path
func Base(path string) string     // returns the last element of path
func Ext(path string) string      // returns the file name extension including the dot
func Abs(path string) (string, error) // returns the absolute path
func Walk(root string, fn WalkFunc) error // walks the file tree (pre-order)
func WalkDir(root string, fn fs.WalkDirFunc) error // like Walk but more efficient; preferred
```

```
p := filepath.Join("music", "pop", "positions.flac")
fmt.Println(filepath.Dir(p)) // music/pop
fmt.Println(filepath.Base(p)) // positions.flac
fmt.Println(filepath.Ext(p)) // .flac
```

`filepath.WalkDir` recursively visits every file in a tree:

```
filepath.WalkDir(".", func(path string, d fs.DirEntry, err error) error {
    if err != nil {
        return err
    }
    if !d.IsDir() {
        fmt.Println(path)
    }
})
```

```
    return nil
}
```



Tip: Use `filepath.Join` instead of string concatenation for paths. It handles leading/trailing separators, double slashes, and platform differences automatically. `filepath.Join("music/", "pop")` returns `"music/pop"`, not `"music//pop"`.

log/slog — Structured Logging

Go's classic `log` package writes plain text lines: great for a quick script, hopeless in production. When your service runs on twenty containers and emits ten thousand lines per minute, you need to search by field, not by regex. **Structured logging** — attaching typed key-value pairs to every log record — makes that possible. Go 1.21 added `log/slog` as the standard structured logger. The Java equivalent is SLF4J + Logback or Log4j2 with a JSON appender, but `slog` is simpler to configure and ships with the standard library.

Handlers and Output Formats

A **handler** controls where and how records are written. `slog` ships two:

```
// logfmt-style text --- readable in a terminal
logger := slog.New(slog.NewTextHandler(os.Stderr, nil))
// time=2026-05-30T15:04:05Z level=INFO msg="track loaded" title=Physical bpm=130

// JSON --- machine-readable; use in production services
logger := slog.New(slog.NewJSONHandler(os.Stderr, nil))
// {"time":"2026-05-30T15:04:05Z","level":"INFO","msg":"track loaded","title":"Physical","bpm":130}
```



Tip: Use `TextHandler` during development (it's readable in a terminal) and `JSONHandler` in production (log aggregators like Datadog, Loki, and Cloud Logging parse it automatically). Switch by reading an environment variable at startup.

Log Levels

`slog` has four built-in levels in increasing severity:

Level	Use for
Debug	Detailed internal state; off in production
Info	Normal events: requests handled, tasks completed
Warn	Recoverable anomalies: retried operations, degraded mode
Error	Failures that need attention but did not crash the process

```
logger.Debug("cache miss", slog.String("key", key))
logger.Info("track loaded", slog.String("title", "Physical"))
logger.Warn("rate limit approaching", slog.Int("remaining", 5))
logger.Error("database query failed", slog.Any("err", err))
```



Trap: Do not log and then return an error for the same event. Either log it (at the point where you handle it) or return it (so the caller can decide whether to log), never both. Doubling up fills logs with duplicate noise and makes it hard to tell where a problem actually originated.

Configuring the Log Level

`slog.HandlerOptions` controls the minimum level and whether source file locations are included:

```
opts := &slog.HandlerOptions{
    Level:    slog.LevelDebug, // log everything
    AddSource: true,          // include file:line in every record
}
logger := slog.New(slog.NewJSONHandler(os.Stderr, opts))
```

For a level you can change at runtime without restarting, use `slog.LevelVar`:

```
var logLevel slog.LevelVar // defaults to Info

func main() {
    if os.Getenv("DEBUG") != "" {
        logLevel.Set(slog.LevelDebug)
    }
    logger := slog.New(slog.NewJSONHandler(os.Stderr, &slog.HandlerOptions{
        Level: &logLevel,
    }))
    slog.SetDefault(logger)
}
```

Setting the Default Logger

`slog.SetDefault(logger)` installs your configured logger as the process-wide default. After this call, the package-level functions `slog.Info(...)`, `slog.Error(...)`, etc. all use your handler. This lets library code call `slog.Info(...)` without knowing which handler the application chose.

```
func main() {
    logger := slog.New(slog.NewJSONHandler(os.Stderr, nil))
    slog.SetDefault(logger) // all subsequent slog.* calls use this logger
    // ...
}
```



Tip: Call `slog.SetDefault` once, early in `main`, before starting any goroutines. Libraries should never call `slog.SetDefault` — that is the application's job.

Typed Attributes

Always use the typed constructors rather than bare strings:

```
func String(key string, value string) Attr // string attribute
func Int(key string, v int) Attr           // int attribute
func Float64(key string, v float64) Attr
func Bool(key string, v bool) Attr
func Duration(key string, v time.Duration) Attr
func Any(key string, value any) Attr       // escape hatch for custom types

logger.Info("track loaded",
    slog.String("title", "Physical"), // Dua Lipa
    slog.Int("bpm", 130),
    slog.Duration("elapsed", time.Since(start)),
)
```



Tip: Keep attribute key names lowercase with underscores (`request_id`, `user_email`, `elapsed_ms`). Consistent naming means you can write the same query in your log aggregator regardless of which service emitted the record. Treat your log schema like an API — it has consumers.

logger.With — Per-Request Context

`logger.With(attrs...)` returns a **child logger** that includes the given attributes on every record it emits. Use it to stamp every log line in a request handler with the request ID and user, so you can filter all lines for one request in one query:

```
func handleUpload(w http.ResponseWriter, r *http.Request) {
    log := slog.Default().With(
        slog.String("request_id", r.Header.Get("X-Request-ID")),
        slog.String("user", userFromContext(r.Context())),
    )
    log.Info("upload started")
    // ... do work ...
    log.Info("upload complete", slog.Int64("bytes", n))
    // both lines carry request_id and user automatically
}
```

This is the Go equivalent of SLF4J's MDC (Mapped Diagnostic Context), but without thread-local storage — the context travels with the logger value, not with the goroutine.

slog.Group — Nested Attributes

`slog.Group` nests attributes under a named key. JSON handlers render it as a nested object; text handlers join the names with a dot:

```
logger.Info("upload complete",
    slog.Group("file",
        slog.String("name", "tracks.csv"),
        slog.Int64("bytes", 204800),
    ),
)
// JSON: {"msg":"upload complete","file":{"name":"tracks.csv","bytes":204800}}
// text: msg="upload complete" file.name=tracks.csv file.bytes=204800
```

Context-Aware Logging

Every log method has a `*Context` variant that accepts a `context.Context` as its first argument:

```
logger.InfoContext(ctx, "query executed", slog.Duration("elapsed", elapsed))
```

Custom handlers can extract trace IDs or span IDs from the context and include them automatically in every record — essential for correlating logs with distributed traces. Even if you do not implement this today, using `*Context` methods now means you can add trace correlation later without changing every log call site.

regexp — Regular Expressions

Go's `regexp` package uses RE2 syntax, which is a strict subset of the PCRE patterns you may know from Java's `java.util.regex`. RE2 guarantees linear-time matching — no catastrophic backtracking — by prohibiting backreferences and look-arounds.

Compile Once, Use Many Times

```
func Compile(expr string) (*Regexp, error) // compile; returns error on bad syntax
func MustCompile(expr string) *Regexp      // compile; panics on bad syntax; use for package-level vars
```

The critical pattern: compile the expression once (at package level or in an init function) and reuse the `*Regexp` value for every match. Compiling on every call is expensive — the equivalent of calling `Pattern.compile(regex)` inside a loop in Java.

```
// At package level: compiled once when the program starts.
var titleRE = regexp.MustCompile(`^[A-Z][a-z]+ \w+$`)
```

```
func isValidTitle(s string) bool {
    return titleRE.MatchString(s)
}
```



Trap: Never call `regexp.MustCompile` (or `regexp.Compile`) inside a function that is called repeatedly. Each call re-parses and re-compiles the pattern — orders of magnitude slower than reusing a compiled `*Regexp`.

Common *Regexp Methods

```
func (re *Regexp) MatchString(s string) bool           // reports whether s matches
func (re *Regexp) FindString(s string) string          // returns leftmost match, or ""
func (re *Regexp) FindAllString(s string, n int) []string // all matches; n=-1 means all
func (re *Regexp) FindStringSubmatch(s string) []string // match + capture groups
func (re *Regexp) ReplaceAllString(src, repl string) string // replace all matches
func (re *Regexp) ReplaceAllStringFunc(src string, f func(string) string) string // replace with function
```

```
var isbnRE = regexp.MustCompile(`\d{3}-\d{10}`)
```

```
fmt.Println(isbnRE.MatchString("978-0135182789")) // true
fmt.Println(isbnRE.FindString("isbn: 978-0135182789 end")) // 978-0135182789
```

```
var versionRE = regexp.MustCompile(`(\d+)\.(\d+)\.(\d+)`)
m := versionRE.FindStringSubmatch("version 1.21.0 released")
fmt.Println(m[0], m[1], m[2], m[3]) // 1.21.0 1 21 0
```

`FindStringSubmatch` returns a slice where `m[0]` is the whole match and `m[1]`, `m[2]`, ... are the capture groups — the same model as Java's `Matcher.group(0)`, `group(1)`, etc.



Tip: Go's RE2 does not support backreferences (`\1`) or lookaheads (`(?=...)`). If you are porting a Java pattern that uses these, you will need to restructure the logic, typically by splitting the match into multiple passes or using `FindStringSubmatch` with post-processing.

cmp and maps — Comparison and Map Utilities

Go 1.21 added the `cmp` and `maps` packages with generic utilities that complement the collections and sorting functions already covered.

cmp — Three-Way Comparison

Go 1.21 added the `cmp` package in the standard library. It provides generic comparison utilities that are useful when sorting slices of structs.


```
// package cmp
func Compare[T Ordered](x, y T) int // -1 if x < y, 0 if x == y, +1 if x > y
```

cmp.Ordered is a type constraint satisfied by all integer types, floating-point types, and string.

The practical use case is sorting a slice of structs by a numeric or string field:

```
import (
    "cmp"
    "fmt"
    "slices"
)

type Track struct {
    Title string
    BPM   int
}

func main() {
    playlist := []Track{
        {Title: "Gata Only",      BPM: 170},
        {Title: "Starboy",        BPM: 186},
        {Title: "Golden Hour",    BPM: 136},
        {Title: "Until I Found You", BPM: 124},
    }

    slices.SortFunc(playlist, func(a, b Track) int {
        return cmp.Compare(a.BPM, b.BPM)
    })

    for _, t := range playlist {
        fmt.Printf("%s: %d\n", t.Title, t.BPM)
    }
}

// Until I Found You: 124
// Golden Hour: 136
// Gata Only: 170
// Starboy: 186
```

cmp.Compare replaces the classic three-way comparison pattern and is safe for floats (it handles NaN consistently).

maps — Map Utilities

Go 1.21 added the maps package to the standard library with utility functions for working with maps.

```
import "maps"

func Keys[Map ~map[K]V, K comparable, V any](m Map) iter.Seq[K] // yields each key
func Values[Map ~map[K]V, K comparable, V any](m Map) iter.Seq[V] // yields each value
func Clone[Map ~map[K]V, K comparable, V any](m Map) Map // shallow copy of m
```

The most common use is combining maps.Keys with slices.Sorted to iterate a map in sorted key order without manually building a key slice:

```
import (
    "fmt"
```

```

    "maps"
    "slices"
)

func main() {
    bpm := map[string]int{
        "amapiano": 112,
        "hyperpop": 160,
        "lo-fi":    85,
    }

    for _, genre := range slices.Sorted(maps.Keys(bpm)) {
        fmt.Printf("%s: %d\n", genre, bpm[genre])
    }
}
// amapiano: 112
// hyperpop: 160
// lo-fi: 85

```

`maps.Clone` makes a shallow copy of a map — useful when you want to mutate a map without affecting the original.

encoding/base64 — Base64 Encoding

`encoding/base64` encodes and decodes binary data as printable ASCII — essential whenever you need to store or transmit raw bytes in a text context (HTTP headers, JSON fields, URLs).

```

func NewEncoder(enc *Encoding, w io.Writer) io.WriteCloser // streaming encoder
func NewDecoder(enc *Encoding, r io.Reader) io.Reader       // streaming decoder
func (enc *Encoding) EncodeToString(src []byte) string      // encode bytes to string
func (enc *Encoding) DecodeString(s string) ([]byte, error) // decode string to bytes

```

The package provides four standard encodings:

Encoding	Use case
<code>base64.StdEncoding</code>	MIME (e-mail, PEM) — uses + and /, padded with =
<code>base64.URLEncoding</code>	URL and filename-safe — uses - and _, padded
<code>base64.RawStdEncoding</code>	Like <code>StdEncoding</code> but no padding
<code>base64.RawURLEncoding</code>	URL-safe, no padding — the most common choice for tokens

```

data := []byte("something in the orange")
encoded := base64.StdEncoding.EncodeToString(data)
fmt.Println(encoded) // c29tZXRoZW5nIGluIHRobzSBvcnFuZ2U=

decoded, err := base64.StdEncoding.DecodeString(encoded)
fmt.Println(string(decoded)) // something in the orange

```



Trap: `StdEncoding` and `URLEncoding` produce different output for the same input. Always use the same encoding for encoding and decoding — a + in `StdEncoding` output becomes %2B in a URL, which a `URLEncoding` decoder will reject.

crypto/rand and math/rand/v2 — Random Numbers

Go has two random number packages for two very different purposes.

crypto/rand produces **cryptographically secure** random bytes suitable for secrets, tokens, and keys. It reads from the OS entropy source (/dev/urandom on Linux, CryptGenRandom on Windows).

```
func Read(b []byte) (n int, err error) // fill b with random bytes; never errors in practice

import "crypto/rand"
import "encoding/base64"

func generateToken(n int) (string, error) {
    b := make([]byte, n)
    if _, err := rand.Read(b); err != nil {
        return "", err
    }
    return base64.RawURLEncoding.EncodeToString(b), nil
}
```

math/rand/v2 (Go 1.22+) is the **fast, non-cryptographic** package for simulations, shuffles, random sampling, and tests. It is not suitable for security-sensitive code.

```
func N[I constraints.Integer](n I) I // uniform random int in [0, n)
func Float64() float64               // uniform float in [0.0, 1.0)
func Shuffle(n int, swap func(i, j int)) // shuffle a slice

import "math/rand/v2"

fmt.Println(rand.N(100)) // random int in [0, 100)
rand.Shuffle(len(tracks), func(i, j int) {
    tracks[i], tracks[j] = tracks[j], tracks[i]
})
```



Trap: Never use math/rand (or math/rand/v2) for passwords, tokens, session IDs, or encryption keys. Use crypto/rand for anything security-sensitive. The fast package is predictable given the seed — a bad actor who knows the seed can predict every value you generate.

crypto/sha256, crypto/aes, crypto/cipher — Hashing and Encryption

Go's crypto subtree provides production-grade primitives. Java programmers familiar with javax.crypto and java.security.MessageDigest will find the same patterns here.

Hashing with crypto/sha256

SHA-256 produces a 32-byte digest — useful for checksums, password hashing (with a proper KDF on top), and HMAC signatures.

```
import "crypto/sha256"

func Sum256(data []byte) [32]byte // one-shot hash
func New() hash.Hash              // streaming hash (implements io.Writer)

digest := sha256.Sum256([]byte("something in the orange"))
fmt.Printf("%x\n", digest) // hex-encoded 64-character string
```



Trap: SHA-256 alone is not suitable for storing passwords. Use golang.org/x/crypto/bcrypt or golang.org/x/crypto/argon2 which are designed to be intentionally slow and include a salt. SHA-256 is fast by design, which makes it easy to brute-force.

Symmetric Encryption with `crypto/aes` and `crypto/cipher`

AES-256-GCM is the standard choice for symmetric encryption in Go: authenticated, fast, and supported by hardware acceleration on most CPUs.

```
import (
    "crypto/aes"
    "crypto/cipher"
    "crypto/rand"
)

func encrypt(key, plaintext []byte) ([]byte, error) {
    block, err := aes.NewCipher(key) // key must be 16, 24, or 32 bytes
    if err != nil {
        return nil, err
    }
    gcm, err := cipher.NewGCM(block) // GCM mode: authenticated encryption
    if err != nil {
        return nil, err
    }
    nonce := make([]byte, gcm.NonceSize()) // 12 bytes; must be unique per message
    if _, err := rand.Read(nonce); err != nil {
        return nil, err
    }
    return gcm.Seal(nonce, nonce, plaintext, nil), nil // nonce prepended to ciphertext
}

func decrypt(key, ciphertext []byte) ([]byte, error) {
    block, _ := aes.NewCipher(key)
    gcm, _ := cipher.NewGCM(block)
    nonce, ct := ciphertext[:gcm.NonceSize()], ciphertext[gcm.NonceSize():]
    return gcm.Open(nil, nonce, ct, nil) // authenticates and decrypts
}
```

The `gcm.Seal` / `gcm.Open` pair handles authentication automatically — if the ciphertext is tampered with, `Open` returns an error. This is called **authenticated encryption with associated data (AEAD)**.



Tip: A 32-byte key gives you AES-256. Generate keys with `crypto/rand` and store them in environment variables or a secrets manager — never hard-code them. The nonce must never be reused with the same key; a fresh random nonce per message guarantees this.

Key Points

- `%+v` adds field names to struct output; `%#v` gives Go-syntax output you can paste back into code.
- `fmt.Fprintf` writes to any `io.Writer` — a file, a buffer, a network connection, a test recorder.
- `io.Copy` streams without buffering the entire input; `io.ReadAll` loads everything into memory.
- `io.TeeReader`, `io.MultiWriter`, `io.LimitReader`, and `io.Pipe` compose readers and writers without intermediate allocations.
- `bufio.Scanner` is the idiomatic line-by-line reader; always check `scanner.Err()` after the loop.

- Always defer `bw.Flush()` after creating a `bufio.Writer` or buffered data will be silently lost.
- Always defer `f.Close()` after opening a file — Go has no try-with-resources.
- `exec.Command` takes separate arguments, not a shell string; use `exec.Command("ls", "-la")`, not `exec.Command("ls -la")`.
- `flag.Parse()` must be called before reading any flag values; call it from `main`, not `init`.
- `time.Duration` is an int64 nanoseconds; write durations as `30 * time.Second`, not `time.Duration(30)`.
- Go's `time.Format` uses reference-time constants (2006-01-02), not pattern tokens like Java's yyyy-MM-dd.
- Always defer `ticker.Stop()` to avoid goroutine leaks.
- `log/slog` (Go 1.21) provides structured logging; use `JSONHandler` in production, `TextHandler` in development.
- Call `slog.SetDefault` once in `main`; use `slog.LevelVar` for a runtime-adjustable level.
- Use `logger.With` to stamp per-request attributes on every log line without repeating them.
- Use `*Context` log methods so custom handlers can inject trace IDs automatically.
- Don't log and return the same error — pick one.
- Compile `regexp.MustCompile` once at package level; never inside a hot function.
- `base64.RawURLEncoding` is the most common choice for URL-safe tokens; always encode and decode with the same variant.
- Use `crypto/rand` for secrets and tokens; use `math/rand/v2` for simulations and shuffles — never the reverse.
- `crypto/sha256` produces a 32-byte digest; use `golang.org/x/crypto/bcrypt` for passwords.
- AES-256-GCM (`crypto/aes + cipher.NewGCM`) provides authenticated encryption; use a fresh random nonce per message.

Exercises

1. **Think about it:** In Java, `InputStream`, `OutputStream`, `Reader`, and `Writer` are four separate abstract class hierarchies. Go has two interfaces — `io.Reader` and `io.Writer` — and a set of composition functions. What design decision makes Go's two-interface model work where Java needed four base classes? What would be harder to express cleanly in Go's model?
2. **What does this print?**

```
package main

import (
    "bufio"
    "fmt"
    "log/slog"
    "os"
    "strings"
    "time"
)

func main() {
    logger := slog.New(slog.NewTextHandler(os.Stdout, &slog.HandlerOptions{
        ReplaceAttr: func(groups []string, a slog.Attr) slog.Attr {
            if a.Key == slog.TimeKey {
                return slog.Attr{} // suppress the timestamp
            }
            return a
        },
    })),
    fmt.Println("Hello, world!")
}
```

```

input := "Physical\nDon't Start Now\nPositions\n"
scanner := bufio.NewScanner(strings.NewReader(input))
count := 0
for scanner.Scan() {
    count++
}

logger.Info("scan complete",
    slog.Int("lines", count),
    slog.Duration("elapsed", 0*time.Millisecond),
)
}

```

3. **Calculation:** You open a 10 MiB file and read it in three ways:

- (a) `os.ReadFile` into a `[]byte`,
- (b) `bufio.NewScanner` reading line by line,
- (c) `io.Copy(io.Discard, f)` using the default 32 KiB copy buffer. For each approach, estimate the peak heap allocation in MiB, assuming the file contains 100,000 lines of 100 bytes each. Which approach is best for counting lines without storing the content?

4. **Where is the bug?**

```

package main

import (
    "fmt"
    "regexp"
)

func countMatches(texts []string, pattern string) int {
    total := 0
    for _, t := range texts {
        re := regexp.MustCompile(pattern)
        if re.MatchString(t) {
            total++
        }
    }
    return total
}

func main() {
    titles := []string{"positions", "Physical", "Don't Start Now", "thank u, next"}
    fmt.Println(countMatches(titles, `^[A-Z]`))
}

```

5. **Write a program:** Write a CLI tool that accepts three flags: `-dir` (a directory path, default `."`), `-ext` (a file extension like `".go"`, default `".go"`), and `-verbose` (a boolean, default `false`). The tool should walk the directory tree, count files whose name ends with the given extension, and print the total. When `-verbose` is set, log each matching file path using `slog` at the `Info` level with a `"file"` attribute. Use `log/slog` with a text handler writing to `os.Stderr`, `path/filepath.WalkDir`, and `flag`.