



Gorgo Go for Java Programmers

May 31, 2026

Contents

18 Testing	1
The testing.T Type	1
Table-Driven Tests	2
t.Helper()	3
Benchmarks	4
Fuzzing	5
Race Detector	5
Goroutine Leak Detection	6
Integration Tests and Build Tags	7
Running Tests	7
Key Points	8
Exercises	8

Chapter 18

Testing

Go's testing package is deliberately minimal: no annotations, no test runner configuration files, no assertion library. What it lacks in ceremony it makes up for in power — table-driven tests, subtests, benchmarks, and a built-in fuzzer all ship with the standard library. If you are used to JUnit 5, most concepts will map cleanly; the idioms are just different.

The `testing.T` Type

Every test function takes a single argument of type `*testing.T`. The naming rule is strict: a test function must be named `TestXxx` where `Xxx` begins with an uppercase letter, and it must live in a file whose name ends in `_test.go`.

```
package music_test

import "testing"

func TestAboutDamnTime(t *testing.T) {
    got := normalize("about damn time")
    want := "About Damn Time"
    if got != want {
        t.Errorf("normalize(%q) = %q, want %q", "about damn time", got, want)
    }
}
```

Go discovers and runs test files automatically with `go test`. There is no `@Test` annotation, no test class — just a naming convention.

`t.Error`, `t.Fatal`, and `t.Log`

The three methods you will reach for most often are:

```
func (t *T) Error(args ...any)           // mark test failed; continue running
func (t *T) Errorf(format string, args ...any) // mark test failed with formatted message; continue running
func (t *T) Fatal(args ...any)           // mark test failed; stop this test immediately
func (t *T) Fatalf(format string, args ...any) // mark test failed with formatted message; stop immediately
func (t *T) Log(args ...any)             // record a message; only printed on failure (or with -v)
func (t *T) Logf(format string, args ...any) // record a formatted message
```

The key difference between `Error` and `Fatal` mirrors the difference between a recoverable and an unrecoverable condition. `t.Error` marks the test as failed but lets it keep running, which is useful when you want to

report multiple independent failures in one pass. `t.Fatal` marks the test as failed and stops the current test function immediately.



Tip: Use `t.Fatal` when subsequent checks depend on a previous one passing. If you set up a database connection in a test and it fails, there is no point running the 20 assertions that follow — use `t.Fatal` to stop early. Use `t.Error` when each check is independent and you want a full picture of all failures.



Trap: `t.Fatal` calls `runtime.Goexit()` under the hood, which unwinds deferred functions in the current goroutine. If you call `t.Fatal` from a goroutine that is not the test's own goroutine, it will **not** stop the test — it will stop only that goroutine. Call `t.Fatal` only from the goroutine that received `t` directly (the test function itself or a helper called from it, not a spawned goroutine).

In Java with JUnit5, you use `Assertions.assertEquals`, `Assertions.assertTrue`, and `Assertions.assertAll`. Go has no assertion library in the standard library. The idiomatic style is a plain `if` that calls `t.Error` or `t.Fatal`. Third-party libraries like github.com/stretchr/testify/assert exist, but many Go teams prefer the standard approach. Whichever you use, your failure messages should state the input, the actual result, and the expected result so that a reader can diagnose the failure without re-running the test. [*test-failure-describes-wrong*]

Table-Driven Tests

Table-driven tests are Go's answer to JUnit 5's `@ParameterizedTest`. Instead of writing one test function per case, you define a slice of structs — each struct is a test case — and loop over them. [*table-driven-tests*]

```
package music_test

import (
    "testing"
)

func TestGoodAsHell(t *testing.T) {
    cases := []struct {
        name  string
        input int
        want  string
    }{
        {name: "zero",    input: 0,   want: "zero stars"},
        {name: "one star", input: 1,   want: "one star"},
        {name: "max",     input: 5,   want: "five stars"},
        {name: "negative", input: -1,  want: "zero stars"},
    }

    for _, tc := range cases {
        t.Run(tc.name, func(t *testing.T) {
            got := starRating(tc.input)
            if got != tc.want {
                t.Errorf("starRating(%d) = %q, want %q", tc.input, got, tc.want)
            }
        })
    }
}
```

The outer test function iterates over the cases and calls `t.Run` for each one. `t.Run` creates a **subtest**: a named, independently tracked test run.

t.Run Subtests

`t.Run(name, func(t *testing.T))` registers and runs a named subtest.

```
func (t *T) Run(name string, f func(t *T)) bool // run f as a named subtest; returns true if f passes
```

Each subtest:

- Has its own pass/fail state.
- Can be run individually with `-run TestFoo/subtest_name`.
- Appears in failure output as `TestFoo/subtest_name`, making it easy to identify which case failed.



Tip: Name your test cases with a name field and pass it as the first argument to `t.Run`. When a case fails, the output shows `--- FAIL: TestGoodAsHell/negative (0.00s)` so you immediately know which row broke. Without subtests you would only see `--- FAIL: TestGoodAsHell` and have to dig through the output to find which input caused it.

In JUnit 5, parameterized tests use `@ParameterizedTest` with `@MethodSource` or `@CsvSource`. Go's table-driven approach with `t.Run` is more explicit — you write a slice literal and a loop — but it gives you the same per-case naming and isolation. Format each `t.Errorf` call with the actual result before the expected result: `got %q, want %q`. [*actual-before-expected*]

t.Helper()

When you extract repeated assertion logic into a helper function, Go's test output points to the **helper** as the failure site rather than the call site in the test. That is rarely what you want. `t.Helper()` fixes this: calling it at the top of a helper function tells the testing framework to attribute any failure in that function to the **caller**.

Without `t.Helper()`:

```
// checkEqual reports whether got == want, but failure points to this line, not the caller.
func checkEqual(t *testing.T, got, want string) {
    if got != want {
        t.Errorf("got %q, want %q", got, want) // file:line points here
    }
}
```

With `t.Helper()`:

```
// checkEqual reports whether got == want.
func checkEqual(t *testing.T, got, want string) {
    t.Helper() // attribute failures to the caller, not this function
    if got != want {
        t.Errorf("got %q, want %q", got, want) // file:line now points to the caller
    }
}
```

Now when `checkEqual` fails, the reported line is the line in your test function that called `checkEqual`, not the line inside `checkEqual` itself.

```
func TestWoman(t *testing.T) {
    checkEqual(t, normalize("woman"), "Woman") // failure reported here
    checkEqual(t, normalize("GOLDEN"), "Golden") // failure reported here
}
```



Tip: Any function that calls `t.Error`, `t.Fatal`, `t.Log`, or another helper should call `t.Helper()` as its first statement. Forgetting `t.Helper` is a common oversight that makes failure output point to the wrong file and line. [*t-helper-for-helpers*]

Benchmarks

A benchmark measures the performance of a piece of code. Benchmarks follow the same file convention as tests — `_test.go` — but the function name starts with `Benchmark` and the argument is `*testing.B`.

```
func BenchmarkNormalize(b *testing.B) {  
    for range b.N {  
        normalize("about damn time")  
    }  
}
```

`b.N` is the number of iterations the framework chooses. The benchmark runner starts with a small `N` and increases it until the total run time is stable enough to report a reliable per-operation time. You do not set `b.N`; the framework sets it for you.

Run benchmarks with `-bench`:

```
go test -bench=. -benchmem ./...
```

The `-benchmem` flag adds allocation counts to the output.

b.ResetTimer

If your benchmark has expensive setup before the measured loop, use `b.ResetTimer()` to exclude the setup time from the measurement.

```
func BenchmarkGolden(b *testing.B) {  
    data := buildLargePlaylist(10_000) // expensive setup  
    b.ResetTimer()                     // start timing only from here  
    for range b.N {  
        _ = findTopTrack(data)  
    }  
}
```

`b.ResetTimer` zeroes the elapsed time and allocation counters so that only the measured loop is included in the reported numbers.



Tip: Always call `b.ResetTimer()` after any setup that you do not want included in the benchmark. Without it, a slow setup inflates the per-operation time, making the benchmark misleading.



Trap: Do not run benchmarks with `go test ./...` by default. Benchmarks only run when you pass `-bench=<pattern>`. Without `-bench`, benchmark functions are compiled but not executed.

In JUnit 5, there is no built-in benchmarking; you need JMH or similar. Go ships a benchmarking harness in the standard library.

Fuzzing

Fuzzing automatically generates inputs to find crashes and panics. A fuzz test is a function named `FuzzXxx` that takes `*testing.F`.

```
func FuzzGoodAsHell(f *testing.F) {
    f.Add("Good as Hell") // seed corpus: start from known inputs
    f.Add("")
    f.Add("ABOUT DAMN TIME")

    f.Fuzz(func(t *testing.T, s string) {
        result := normalize(s)
        if len(result) > 0 && result[0] >= 'a' && result[0] <= 'z' {
            t.Errorf("normalize(%q) starts with lowercase: %q", s, result)
        }
    })
}
```

`f.Add` seeds the fuzzer with known inputs. The fuzzer uses those seeds as a starting point and then mutates them to generate new inputs.

`f.Fuzz` registers the function that is called for each generated input. The first argument is always `*testing.T`; the remaining arguments match the types passed to `f.Add`.

Run fuzzing with `-fuzz`:

```
go test -fuzz=FuzzGoodAsHell -fuzztime=30s ./...
```

Without `-fuzz`, a fuzz test runs only the seed corpus — just like a regular test. With `-fuzz`, the engine runs indefinitely (or until `-fuzztime` elapses) mutating inputs until it finds a failure. When a failure is found, the engine writes the failing input to `testdata/fuzz/FuzzXxx/` so you can reproduce it.



Tip: Fuzz tests double as regression tests. The seed corpus (`f.Add` calls) runs every time `go test` runs — no `-fuzz` flag needed. Add the `testdata/fuzz/` directory to version control so that previously found failures are always checked.



Wut: Fuzzing is not available as a built-in in JUnit 5; you need a separate library. Go 1.18 added native fuzzing to the standard toolchain.

Race Detector

Chapter 11 introduced the race detector briefly. Here is how to use it in tests.

Run `go test -race` to enable the race detector:

```
go test -race ./...
```

The race detector instruments the binary to track every memory access. If two goroutines access the same variable concurrently and at least one of them is a write, the detector prints a detailed report showing the goroutine stacks at both access sites.



Tip: Run `go test -race` in CI on every push. The race detector has a small runtime cost (typically 2–20x slowdown), which is acceptable in CI but may be too slow for production. The cost of finding a race in production is far higher than the cost of running a slower test suite.



Trap: The race detector only catches races that **actually occur at runtime**. It cannot prove the absence of races. A test suite with low concurrency coverage might miss a race that the detector would catch under higher load. Write tests that exercise concurrent code paths.

Here is a concise example of a race the detector will catch:

```
func TestCounterRace(t *testing.T) {
    var count int
    var wg sync.WaitGroup

    for i := 0; i < 100; i++ {
        wg.Add(1)
        go func() {
            defer wg.Done()
            count++ // data race: concurrent unsynchronized write
        }()
    }
    wg.Wait()
    t.Log("count:", count)
}
```

Run this with `go test -race` and the detector reports the race immediately.

Goroutine Leak Detection

Chapter 12 covered goroutine leaks in depth: what causes them, how to prevent them with context cancellation and done channels, and how `go.uber.org/goleak` detects them in tests. Every goroutine you spawn should have a clear exit path; if it has none, it is a leak. [*goroutine-must-exit*] The testing-specific summary: place `goleak.VerifyTestMain(m)` in `TestMain` to check every test in the package, or `defer goleak.VerifyNone(t)` in individual tests.

```
// Check all tests in the package --- preferred.
func TestMain(m *testing.M) {
    goleak.VerifyTestMain(m)
}

// Check a single test --- use when you cannot modify TestMain.
func TestAboutDamnTimeStream(t *testing.T) {
    defer goleak.VerifyNone(t)
    // ...
}
```



Tip: `TestMain` is Go's equivalent of JUnit 5's `@BeforeAll` / `@AfterAll` at the package level. It runs once per test binary rather than once per test function — the right place for global setup such as opening a shared database connection or installing `goleak`.



Wut: `goleak.VerifyTestMain` calls `m.Run()` internally. Do not call `m.Run()` yourself before passing `m` to it — the tests would run twice.

Integration Tests and Build Tags

Unit tests run quickly and need no external dependencies. Integration tests talk to real databases, real HTTP servers, or real message queues — they are slower and require infrastructure. Build tags let you separate the two so that `go test ./...` runs only the fast unit tests by default.

A build tag is a comment at the very top of a file, before the package declaration:

```
//go:build integration
```

```
package music_test
```

This file is excluded from the build unless the `integration` tag is provided. To run integration tests:

```
go test -tags=integration ./...
```

To run unit tests only:

```
go test ./...
```



Tip: Use separate CI pipeline stages: one stage runs `go test ./...` on every commit (fast, no infrastructure), and a second stage runs `go test -tags=integration ./...` before merging to main (slower, requires a test database or test container).



Wut: In Go 1.17 and earlier, build tags used a different syntax: `// +build integration` (a comment, not a `//go:build` directive). Go 1.17 introduced the `//go:build` form, which is now the standard. `gofmt` will add the `//go:build` form for you if you only write the old form. Both can coexist in the same file for backward compatibility.

Running Tests

The `go test` command is the entry point for all test-related tasks.

```
go test ./...
```

Run all tests in the current module:

```
go test ./...
```

The `./...` pattern matches the current directory and all subdirectories.

-count=1 (Disable Caching)

Go caches test results. If the test sources and dependencies have not changed, `go test` reuses the cached result rather than re-running the tests. This is usually helpful, but sometimes you want to force a fresh run — for example, when testing code that depends on external state.

```
go test -count=1 ./...
```

`-count=1` tells the test runner to run each test exactly once and bypass the cache.



Tip: Use `-count=1` in CI to guarantee that tests always run, even when nothing has changed in the source. Cached test results in CI can hide flaky tests.

-timeout

By default `go test` applies a 10-minute timeout to the entire test binary. You can override this:

```
go test -timeout=30s ./...
```

Set a tight timeout in CI so that a hanging test (for example, a goroutine waiting on a channel that is never closed) fails fast rather than blocking your pipeline for 10 minutes.

-run and -bench

`-run` filters which test functions run by matching against a regular expression. `-bench` does the same for benchmarks.

```
go test -run=TestGoodAsHell ./...           # run only tests matching "TestGoodAsHell"
go test -run=TestGoodAsHell/zero ./...      # run only the "zero" subtest
go test -bench=BenchmarkNormalize ./...     # run only this benchmark
go test -bench=. -benchmem ./...           # run all benchmarks with allocation stats
```

A Typical CI Invocation

A minimal, correct CI test command:

```
go test -race -count=1 -timeout=120s ./...
```

This runs all tests with the race detector, bypasses the cache, and fails if anything takes more than two minutes.

Key Points

- Test functions are named `TestXxx` and live in `_test.go` files; no annotations required.
- `t.Error` marks a failure and continues; `t.Fatal` marks a failure and stops the test immediately.
- Table-driven tests use a slice of anonymous structs; `t.Run` creates named subtests for each case.
- `t.Helper()` must be called at the top of any helper function so that failure output points to the call site, not the helper.
- Benchmarks are named `BenchmarkXxx` and receive `*testing.B`; `b.N` is the iteration count set by the framework.
- Call `b.ResetTimer()` after setup to exclude setup time from benchmark measurements.
- Fuzz tests are named `FuzzXxx` and receive `*testing.F`; `f.Add` seeds the corpus, `f.Fuzz` registers the test body.
- Without `-fuzz`, a fuzz test runs only the seed corpus, acting as a standard test.
- Run `go test -race` in CI on every push to catch data races; the race detector only reports races that actually occur.
- `goleak.VerifyTestMain(m)` in `TestMain` detects goroutine leaks after the test suite finishes.
- Use `//go:build integration` to gate slow integration tests; run them with `-tags=integration`.
- `go test ./...` runs all tests; `-count=1` disables caching; `-timeout` caps the run time.

Exercises

1. **Think about it:** JUnit 5's `@ParameterizedTest` with `@CsvSource` and Go's table-driven tests with `t.Run` both let you run the same logic against many inputs. Describe two concrete advantages that Go's table-driven approach gives you over `@CsvSource`. Then explain the key behavioral difference between `t.Fatal` and `t.Error` inside a subtest, and describe a scenario where you would deliberately choose `t.Error` over `t.Fatal`.
2. **What does this print?** Trace the output when this test is run with `go test -v`.

```

package music_test

import "testing"

func checkPositive(t *testing.T, n int) {
    if n <= 0 {
        t.Errorf("expected positive, got %d", n)
    }
}

func TestGolden(t *testing.T) {
    checkPositive(t, 1)
    t.Log("checked 1")
    checkPositive(t, -1)
    t.Log("checked -1")
    checkPositive(t, 2)
    t.Log("checked 2")
}

```

3. **Calculation:** A benchmark function has the following structure:

```

func BenchmarkWoman(b *testing.B) {
    for range b.N {
        _ = processTrack("Woman")
    }
}

```

On the first probe the framework sets `b.N = 1` and measures elapsed time. It then sets `b.N = 100`, then `b.N = 10_000`, then `b.N = 1_000_000`. The framework stops when the total elapsed time exceeds one second. If `processTrack` takes exactly 500 ns per call, at which value of `b.N` does the total elapsed time first exceed one second? What is the reported ns/op value?

4. **Where is the bug?** The following test helper is supposed to make failure output point to the call site in `TestAboutDamnTime`, but it does not. Identify the bug and show the fix.

```

package music_test

import "testing"

func assertNormalized(t *testing.T, input, want string) {
    got := normalize(input)
    if got != want {
        t.Fatalf("normalize(%q): got %q, want %q", input, got, want)
    }
}

func TestAboutDamnTime(t *testing.T) {
    assertNormalized(t, "about damn time", "About Damn Time")
    assertNormalized(t, "good as hell", "Good as Hell")
}

```

5. **Write a program:** Write a table-driven test for the following function. Your test must include at least five cases covering normal input, empty string, all-caps input, and a multi-word title. Use `t.Run` for each case and `t.Helper` in any helper you write.

```

// TitleCase converts a string to title case.
// Each word's first letter is uppercased; the rest are lowercased.

```

```
// Words are separated by spaces.  
func TitleCase(s string) string
```